

Autonomous Articulated Vehicle Control via Deep Reinforcement Learning and High-Dimensional Perception

Benjamin Schnapp

March 30, 2026

Contents

1	Introduction	5
1.1	Motivation	5
1.2	Problem Definition	5
1.3	Thesis Contributions	6
1.4	Institutional Context	6
1.5	Thesis Organisation	7
2	Literature Review	8
2.1	Articulated Vehicle Control	8
2.1.1	Classical Stabilisation Techniques	8
2.1.2	Kinematic Models	8
2.2	Reinforcement Learning in Autonomous Driving	9
2.2.1	Policy Gradient Methods	9
2.2.2	Transformer-Based Methods	9
2.3	Perception and Representation Learning	9
2.3.1	Convolutional Neural Networks for Driving	9
2.3.2	Autoencoders for State Compression	9
2.4	Safety in Autonomous Systems	10
2.5	Benchmarking Against Recent Waterloo Theses	10
3	Mathematical Modeling and Problem Formulation	12
3.1	Tractor Dynamics	12
3.2	Trailer Articulation and Reverse Instability	12
3.3	Challenges of Reverse Maneuvering	13
3.4	MDP Formulation	13
4	Perception and Observation Space Design	14
4.1	Bird's-Eye View Observation	14
4.2	CNN Encoder	14
4.3	Autoencoder Latent Space	15
4.4	Observation Space Summary	15

5	Reinforcement Learning Framework	16
5.1	Twin Delayed Deep Deterministic Policy Gradient (TD3)	16
5.1.1	Reward Function	16
5.1.2	Multiplicative vs. Additive Reward Structures	17
5.2	Decision Transformer	17
5.2.1	Dataset Requirements	17
5.3	Hyperparameter Optimisation	18
6	Simulation Environment	19
6.1	Custom Pygame/Gymnasium Environment	19
6.2	Environment Tasks	20
6.2.1	Line Following	20
6.2.2	Obstacle Avoidance	20
6.3	Vehicle Parameters	20
6.4	Action Space	21
6.5	Testing Protocols	21
6.5.1	Lane Following Performance	21
6.5.2	Obstacle Avoidance	21
6.5.3	Comparison with Prior Work	21
7	Results and Discussion	23
7.1	Path Following Performance	23
7.2	Reverse Maneuvering Success	23
7.3	Perception Robustness	23
7.4	Benchmark Comparison	23
8	Conclusion	25
8.1	Summary of Findings	25
8.2	Limitations	25
8.3	Future Work	26
A	Simulation Software Architecture	27
B	Full Hyperparameter Tables	28
C	Additional Plots	29
D	Reward Function Derivation	30

List of Figures

List of Tables

1.1	MASc thesis requirements summary	7
2.1	Selected recent MASc theses, University of Waterloo Mechatronics Engineering	10
4.1	Observation space components	15
6.1	Vehicle parameters used in simulation	20
6.2	Lane following experiment definitions	21
7.1	Controller performance comparison	23
7.2	Experiment matrix: tasks, observations, and success criteria	24
B.1	TD3 Hyperparameters	28

Chapter 1

Introduction

1.1 Motivation

The increasing demand for autonomous logistics and the growth of large-scale goods transportation have positioned articulated tractor-trailer vehicles as a critical target for autonomous control research. The transition from rigid-body autonomous navigation to the control of articulated tractor-trailer systems represents a significant escalation in mechanical complexity, perception requirements, and control instability, particularly during reverse maneuvering.

Unlike single-body vehicles, tractor-trailer systems introduce non-holonomic constraints at the hitch point that fundamentally alter the stability characteristics of the vehicle. In reverse motion, the articulation angle γ between the tractor and trailer exhibits open-loop instability: absent precise closed-loop control, any perturbation grows unbounded, culminating in the well-known jackknife condition from which no steering correction can recover the trailer’s alignment. Classical approaches such as Nonlinear Model Predictive Control (NMPC) can manage this instability by optimising over a prediction horizon, but their computational demands and reliance on accurate system models limit their applicability in unstructured or sensor-rich environments.

1.2 Problem Definition

The core challenge addressed in this thesis is twofold. First, the kinematic and dynamic complexity of tractor-trailer systems makes it difficult to design reward functions and control policies that simultaneously optimise path tracking performance and hitch stability. Second, real-world autonomous deployment requires perception of lane boundaries and obstacles directly from high-dimensional sensor data, rather than from pre-computed symbolic state representations.

Prior work in the ME 780 project evaluated Deep Deterministic Policy Gradient

(DDPG) and Decision Transformer (DT) models for path following in single-vehicle systems. While DDPG achieved cross-track errors competitive with PID control, both learning-based methods showed significant degradation under novel conditions, and neither was tested on articulated dynamics or reverse motion.

1.3 Thesis Contributions

This thesis extends the ME 780 and ECE 682 foundational work by:

- developing a full kinematic and dynamic model of a tractor-trailer system including hitch instability analysis, building on the bicycle model and Pacejka tire dynamics established in prior work;
- designing a vision-based RL observation space using CNN encoders and autoencoders operating on bird’s-eye-view (BEV) images rendered from a custom Pygame simulation environment;
- engineering a reward function that balances cross-track error, heading error, hitch articulation stability, and collision avoidance;
- implementing and benchmarking Twin Delayed DDPG (TD3) agents in lane-following and obstacle-avoidance tasks within a custom Gymnasium-based simulation environment;
- comparing TD3 performance with classical control baselines (PID, MPC) and the prior ME 780 DDPG results across path-following metrics.

The research aligns with the objectives of the Mechatronics Vehicle Systems (MVS) Lab and the WATonoBus project at the University of Waterloo.

1.4 Institutional Context

This work is submitted in partial fulfilment of the requirements for a Master of Applied Science (MASc) degree in the Department of Mechanical and Mechatronics Engineering at the University of Waterloo. The degree requires six terms of full-time registration, completion of 3–5 graduate-level courses, appointment of two additional faculty readers, and a mandatory 15-day public display period. The target completion date is early July, with a seminar scheduled for late July.

Table 1.1: MASc thesis requirements summary

Element	Requirement
Research Scope	Substantial problem extending beyond existing laboratory results
Literature Review	Comprehensive assessment of classical control and modern AI approaches
Validation	Simulation-based empirical evidence across multiple environment configurations
Length/Format	Typically 80–120 pages; GSPA formatting and UWSpace submission
Committee	Supervisor plus two faculty readers; formal display period

1.5 Thesis Organisation

The remainder of this thesis is structured as follows. Chapter 2 reviews classical and learning-based control methods for articulated vehicles, alongside relevant perception and safety literature. Chapter 3 derives the state-space equations for the tractor-trailer system and analyses the jackknife instability. Chapter 4 presents the CNN and autoencoder-based perception stack. Chapter 5 defines the MDP formulation and details the DDPG and Decision Transformer implementations. Chapter 6 describes the custom Pygame/-Gymnasium simulation environment and experimental setup. Chapter 7 presents and discusses experimental results. Chapter 8 summarises the findings and outlines future work.

Chapter 2

Literature Review

2.1 Articulated Vehicle Control

2.1.1 Classical Stabilisation Techniques

Classical articulated vehicle control methods address the hitch instability through structured model-based approaches:

- **PID control:** Simple to implement and computationally negligible, but incapable of anticipating the jackknife condition. Gains must be re-tuned for different payloads and path curvatures.
- **Model Predictive Control (MPC):** Optimises steering over a receding horizon subject to state and input constraints. Linear MPC can enforce hard bounds on the articulation angle γ .
- **Nonlinear MPC (NMPC):** Captures the full nonlinear kinematics including reverse instability. Optimises steering over a receding horizon subject to state and input constraints, enabling hard bounds on the articulation angle γ , but with higher computational cost than linear methods.
- **Active hitch stabilisation:** Applies an additional actuated degree of freedom at the hitch to damp articulation oscillations. Research at Waterloo into trailer sway control and active hitches demonstrates that lateral articulation at the hitch can mitigate oscillations, but for a standard autonomous tractor-trailer this must be managed purely through the tractor's steering angle δ .

2.1.2 Kinematic Models

The bicycle model provides a baseline for rigid-body longitudinal and lateral dynamics. For an articulated system, the state vector must be expanded to include the trailer heading and articulation angle, as detailed in Chapter 3.

2.2 Reinforcement Learning in Autonomous Driving

2.2.1 Policy Gradient Methods

Deep Deterministic Policy Gradient (DDPG) is an off-policy actor-critic algorithm designed for continuous action spaces. Prior work in ME 780 established that DDPG achieves a cross-track error of approximately 0.505 m for rigid-body path following, outperforming a PID baseline of 0.784 m on the same task.

Proximal Policy Optimisation (PPO) offers improved stability over DDPG through clipped probability ratios, but its on-policy sample efficiency is lower, which is a significant consideration when simulation rollouts are expensive.

2.2.2 Transformer-Based Methods

Decision Transformers (DT) reformulate RL as sequence modelling: an action is predicted from a context window of past states, actions, and returns-to-go using a causal self-attention mechanism. The ME 780 study found that DTs achieve a cross-track error of 1.717 m and are prone to oscillation if the training dataset is insufficiently diverse. To address this, the MASc research must utilise a dataset that includes recovery maneuvers, either by injecting noise into a DDPG expert’s actions or by explicitly including high-error-state trajectories.

2.3 Perception and Representation Learning

2.3.1 Convolutional Neural Networks for Driving

CNNs have become the standard front-end for autonomous driving perception, extracting spatial features from camera images and occupancy grids. End-to-end learning approaches train CNNs directly from sensor inputs to control outputs, allowing the network to learn which visual features are most relevant for the control task. In the simulation setting of this thesis, CNNs process bird’s-eye view images rendered from the Pygame environment to extract lane geometry and obstacle information.

2.3.2 Autoencoders for State Compression

Autoencoders compress high-dimensional observations into a compact latent representation, mitigating the curse of dimensionality for RL training. By pre-training an autoencoder on images collected from simulation rollouts, the encoder learns to capture essential geometric features of the road and obstacles. The RL agent then operates on the latent vector z , which reduces policy network parameters, accelerates convergence, and improves generalisation across environments.

2.4 Safety in Autonomous Systems

Safety of the Intended Functionality (SOTIF) defines the safety boundary for autonomous systems operating under nominal sensor and actuator conditions. For articulated vehicles, a key safety concern is the jackknife condition: the RL reward function must penalise large articulation angles to discourage the agent from entering unstable configurations. Control Barrier Functions (CBFs) offer a formal mechanism to overlay hard safety constraints on top of a learned policy, and are identified as a direction for future work beyond the simulation-based experiments of this thesis.

2.5 Benchmarking Against Recent Waterloo Theses

To contextualise the expected level of rigor, Table 2.1 summarises recent MASc theses from the Mechatronics domain at the University of Waterloo.

Table 2.1: Selected recent MASc theses, University of Waterloo Mechatronics Engineering

Author (Year)	Topic	Methodology	Rigor Highlights
Joseph (2025)	Gaze-enabled grasping assistance	ROS2, Kinova Arm, Meta Quest Pro	User study with 30 participants; workload and performance metrics
Thibault (2022)	Humanoid robotic manipulation	EUROBENCH, REEM-C robot	Defined “loco-manipulation”; new manipulability–stability metrics
Hart (2023)	Gait detection via machine learning	Wearable sensors, ternary classification	6 classical + 1 Transformer model; grid-search hyperparameter optimisation
Wei (2021)	Path following for manipulators	Transpose Jacobian control	Free of inverse transformations; validated on 2-DoF and 4-DoF robots
Bickford (2019)	Autonomous bicycle stabilisation	Linear control, Whipple model	Multi-loop architecture; “virtual paths” for high-error recovery

Common themes include strong mathematical modelling foundations and comprehensive validation datasets. For this thesis, a simple “it works” result is insufficient; the

candidate must quantify stability limits, evaluate performance across varied paths, and compare against classical benchmarks.

Chapter 3

Mathematical Modeling and Problem Formulation

3.1 Tractor Dynamics

The tractor longitudinal and lateral dynamics are:

$$\dot{x} = v_x \cos \psi - v_y \sin \psi \quad (3.1)$$

$$\dot{y} = v_x \sin \psi + v_y \cos \psi \quad (3.2)$$

where v_x is the longitudinal velocity, v_y is the lateral velocity, and ψ is the vehicle heading.

3.2 Trailer Articulation and Reverse Instability

The articulated system introduces hitch geometry defined by:

- L_h : distance from tractor rear axle to hitch point
- L_t : distance from hitch point to trailer axle
- γ : articulation angle between tractor and trailer

The hitch angle kinematics during reverse motion are:

$$\dot{\gamma} = \frac{v}{L_t} \sin \gamma - \frac{v}{L_h} \tan \delta \cos \gamma \quad (3.3)$$

Reverse motion introduces instability as γ approaches the jackknife threshold.

3.3 Challenges of Reverse Maneuvering

Reverse motion is characterised by a sign reversal in the kinematic equations, which transforms the stabilising effect of forward velocity into a destabilising force. Successful reverse maneuvering requires anticipating the steering input needed to prevent γ from reaching its mechanical limits. The RL agent developed in this thesis must learn analogous foresight: rewarding actions that keep the trailer aligned with the target path while maintaining γ within a safe operational envelope.

3.4 MDP Formulation

The control problem is cast as a Markov Decision Process (MDP) $(\mathcal{S}, \mathcal{A}, \mathcal{T}, \mathcal{R}, \gamma_d)$ where:

- $\mathcal{S}$: observation space consisting of a vector of state variables (steering angle, articulation angle, cross-track error, heading error) and an 84×84 grayscale BEV image;
- $\mathcal{A}$: continuous action space $[\dot{\delta}, v] \in [-15^\circ/\text{s}, 15^\circ/\text{s}] \times [0, 2 \text{ m/s}]$;
- $\mathcal{T}$: transition dynamics simulated in the custom Pygame/Gymnasium environment;
- $\mathcal{R}$: reward function defined in Chapter 5;
- γ_d : discount factor.

Episodes terminate on collision, jackknife ($|\gamma| > 90^\circ$), or successful path completion.

Chapter 4

Perception and Observation Space Design

4.1 Bird’s-Eye View Observation

The RL agent receives observations from two complementary sources: a vector of state variables and a bird’s-eye view (BEV) image rendered from the simulation environment. This observation design is motivated by the prior ME 780 work, which used only vector state observations (cross-track error, heading error, velocity), and extends it to include visual perception of lane boundaries and obstacles.

The BEV image is generated by a configurable camera anchored to the tractor’s rear axle, rendering a top-down 84×84 grayscale image of the vehicle and its surroundings from the Pygame simulation. The anchor point and zoom level are configurable in the simulation settings. The image encodes lane markings, obstacle positions, and the vehicle geometry relative to the lane, providing the agent with spatial context that the vector observation alone cannot capture.

4.2 CNN Encoder

A Convolutional Neural Network (CNN) serves as the front-end of the RL agent, processing the BEV image to extract spatial features relevant to lane tracking and obstacle avoidance. The CNN enables the agent to perceive the environment directly from rendered images, without requiring a pre-defined symbolic representation of every object.

The CNN architecture is implemented as a custom feature extractor within the Stable Baselines3 framework, following an ImageEncoder design. The extracted image features are concatenated with the vector state observations before being passed to the TD3 actor and critic networks. This combined observation allows the policy to leverage both precise state information and rich spatial context.

The CNN is trained end-to-end with the TD3 agent, allowing the feature extractor to learn which visual features are most relevant for path maintenance and collision avoidance.

4.3 Autoencoder Latent Space

To improve RL training efficiency, an autoencoder can be used to compress high-dimensional BEV image observations into a compact latent representation. The encoder $E(\cdot)$ maps the image observation s to a latent vector z , and the decoder $D(\cdot)$ reconstructs the original image:

$$z = E(s) \quad (4.1)$$

$$\hat{s} = D(z) \quad (4.2)$$

The autoencoder is pre-trained on images collected from simulation rollouts, allowing the encoder to learn compact representations of the road geometry and obstacle configurations. The RL agent then operates on z rather than the raw image, reducing the number of parameters the policy network must learn and potentially improving convergence.

4.4 Observation Space Summary

The full observation space is a dictionary with two components:

Table 4.1: Observation space components

Component	Description	Role
vector	State variables: steering angle, hitch angle, cross-track error, heading error	Provides precise numerical state information to the policy
image	84×84 grayscale BEV image rendered from Pygame simulation	Encodes spatial context: lane markings, obstacle positions, vehicle geometry

Collision detection is performed using Pygame sprite mask overlap, enabling pixel-precise detection of contact between the vehicle body and obstacles or lane boundaries. Episode termination is triggered on collision, jackknife ($|\gamma| > 90^\circ$), or path completion.

Chapter 5

Reinforcement Learning Framework

5.1 Twin Delayed Deep Deterministic Policy Gradient (TD3)

Twin Delayed Deep Deterministic Policy Gradient (TD3) is an off-policy actor-critic algorithm designed for continuous action spaces. TD3 extends the Deep Deterministic Policy Gradient (DDPG) algorithm from the ME 780 prior work with three key improvements: twin critic networks (to reduce overestimation bias), delayed policy updates (to improve stability), and target policy smoothing (to reduce variance in critic updates).

Like DDPG, the TD3 actor network $\mu(s \mid \theta^\mu)$ maps states to deterministic actions, while the critic networks $Q_1, Q_2(s, a \mid \theta^{Q_1}, \theta^{Q_2})$ evaluate action quality. The target value is computed using the minimum of the two critic estimates:

$$y = r_t + \gamma_d \min_{i=1,2} Q_i(s_{t+1}, \mu(s_{t+1})) + \epsilon \quad (5.1)$$

where $\epsilon \sim \text{clip}(\mathcal{N}(0, \sigma), -c, c)$ is clipped Gaussian noise added to the target action for smoothing. Parameters are updated by minimising the temporal difference error in the critics and maximising the expected return through the actor using the deterministic policy gradient theorem.

TD3 is implemented using the Stable Baselines3 library, which provides a reliable implementation integrated with the Gymnasium environment API.

5.1.1 Reward Function

The proposed reward function is:

$$R = w_1 e^{-|e_y|} + w_2 e^{-|e_\psi|} - w_3 |\gamma| - w_4 C \quad (5.2)$$

where:

- e_y = trailer cross-track error
- e_ψ = heading error
- γ = articulation angle
- C = collision penalty

A multiplicative alternative is also considered:

$$R = e^{-|e_y|} \cdot e^{-|e_\psi|} \quad (5.3)$$

5.1.2 Multiplicative vs. Additive Reward Structures

Prior ME 780 results indicated that a multiplicative reward structure forces the agent to minimise both cross-track and heading errors simultaneously, preventing the policy from exploiting partial credit in the additive form. In this thesis, this logic is extended to include hitch articulation stability: the multiplicative penalty on $|\gamma|$ discourages the agent from trading path accuracy for a large articulation angle.

5.2 Decision Transformer

The Decision Transformer (DT) reformulates RL as a sequence-to-sequence problem. Given a context window of past states $s_{t-k:t}$, actions $a_{t-k:t-1}$, and returns-to-go $\hat{R}_{t-k:t}$, the model predicts the next action using the causal self-attention mechanism of the Transformer architecture:

$$a_t = \pi_\theta(s_{t-k:t}, a_{t-k:t-1}, \hat{R}_{t-k:t}) \quad (5.4)$$

Decision transformers are an offline reinforcement learning approach: rather than training by interacting with the environment, they are trained on recorded trajectories of an agent (such as the TD3 policy) interacting with the environment. This architecture is relevant for the tractor-trailer setting, where early steering errors may only result in a jackknife condition several seconds later, requiring the model to reason over long-term dependencies.

5.2.1 Dataset Requirements

The ME 780 study revealed that DTs are prone to oscillation if the training dataset is too narrow in state diversity. To address this, training data should include recovery manoeuvres obtained by:

- injecting noise into a TD3 expert’s actions during data collection;

- explicitly including trajectories initialised from high-error states (e.g., $|e_y| > 1\text{ m}$ or $|\gamma| > 20^\circ$);
- including step discontinuities in the reference path to force the model to learn re-acquisition behaviour, as demonstrated in the ME 780 work.

5.3 Hyperparameter Optimisation

Both the TD3 and DT agents require careful hyperparameter tuning. The actor-critic learning rates, replay buffer size, target network update rate, and noise scale are selected via grid search over simulation rollouts. Full hyperparameter tables are provided in Appendix [B](#).

Chapter 6

Simulation Environment

6.1 Custom Pygame/Gymnasium Environment

The training and evaluation environment is a custom 2D simulation built on the Farama Foundation Gymnasium library, consistent with the approach taken in the ME 780 prior work. Gymnasium provides a standardised API for implementing environments that integrate directly with reinforcement learning algorithm libraries, including Stable Baselines3.

The simulation is implemented using Pygame for rendering and physics integration. At each timestep, the discrete-time tractor-trailer state-space equations are re-evaluated using the previous state and control inputs, updating the vehicle position, heading, hitch angle, and associated error signals. Pygame renders the scene as a 2D top-down view, from which the BEV image observation is captured.

The simulation environment includes:

1. **Vehicle Model:** A combined dynamic tractor (bicycle model with Pacejka tire dynamics) and kinematic trailer with hitch constraints, using Tesla Model S parameters as in the ME 780 and ECE 682 prior work.
2. **Lane Generation:** Randomly generated lane paths using cubic spline interpolation, producing varying curvature paths for training and evaluation.
3. **Obstacle Placement:** Randomly placed static obstacles within the lane, used in the obstacle avoidance task.
4. **Collision Detection:** Pygame sprite mask-based collision detection, providing pixel-precise contact sensing between the vehicle body and obstacles or lane boundaries.
5. **BEV Camera:** A configurable bird’s-eye view camera anchored to the tractor’s rear axle, rendering an 84×84 grayscale image of the vehicle and its surroundings.

6.2 Environment Tasks

Two distinct Gymnasium environments are implemented, corresponding to the two primary research tasks:

6.2.1 Line Following

The lane-following environment extends the base tractor-trailer environment to present a lane-tracking task. A reference lane is generated at the start of each episode using cubic spline interpolation through randomly placed waypoints. The reward function encourages the agent to keep both the tractor and trailer within the lane boundaries while maintaining forward progress.

The observation at each timestep includes:

- **vector**: steering angle δ , hitch articulation angle γ , cross-track error e_y , heading error e_ψ , and their time derivatives;
- **image**: an 84×84 grayscale BEV image rendered from the tractor rear axle anchor point.

6.2.2 Obstacle Avoidance

The obstacle avoidance environment adds static obstacles within the lane. The agent must navigate the tractor-trailer through the lane while avoiding obstacles, which requires both path-following behaviour and collision avoidance. Episode termination occurs on collision, jackknife ($|\gamma| > 90^\circ$), or successful path completion.

6.3 Vehicle Parameters

The simulation uses the same vehicle parameters as the ME 780 and ECE 682 prior work, based on Tesla Model S specifications:

Table 6.1: Vehicle parameters used in simulation

Parameter	Value	Parameter	Value
m	1500 kg	l_f	1.2 m
I_z	3000 kg m ²	l_r	1.6 m
C_f	80000 N/rad	C_d	0.208
C_r	80000 N/rad	A	2.4 m ²
Δt	0.1 s	ρ	1.225 kg/m ³

6.4 Action Space

The agent controls the vehicle through a continuous two-dimensional action space:

$$\mathbf{a} = [\dot{\delta}, v] \in [-15^\circ/\text{s}, 15^\circ/\text{s}] \times [0, 2\text{ m/s}] \quad (6.1)$$

where $\dot{\delta}$ is the steering rate (rate of change of steering angle) and v is the commanded speed. This action parameterisation is consistent with the observation from the ECE 682 work that the steering angle cannot be set instantaneously on a physical vehicle; controlling the rate of change is more physically realistic.

6.5 Testing Protocols

6.5.1 Lane Following Performance

Table 6.2: Lane following experiment definitions

Experiment	Objective	Success Metric
Path Fidelity (Tractor)	Maintain low CTE for tractor on varying curvature paths	Mean Absolute CTE
Path Fidelity (Trailer)	Maintain low CTE for trailer, ensuring full vehicle stays in lane	Mean Absolute CTE
Hitch Stability	Keep articulation angle within safe bounds	Max $ \gamma $ per episode

6.5.2 Obstacle Avoidance

The obstacle avoidance task tests the agent’s ability to navigate through a lane containing static obstacles. Success is measured by collision rate and the agent’s ability to complete the lane while maintaining hitch stability.

6.5.3 Comparison with Prior Work

Results are compared against:

- the DDPG model from ME 780 (cross-track error 0.505 m on single-vehicle path following);
- a PID baseline from ME 780 (cross-track error 0.784 m);

- an LQR-based baseline from ECE 682;
- an MPC controller implemented using OSQP.

Chapter 7

Results and Discussion

7.1 Path Following Performance

7.2 Reverse Maneuvering Success

7.3 Perception Robustness

7.4 Benchmark Comparison

Table 7.1: Controller performance comparison

Model	Avg. CTE (m)	Jackknife Events/100	Latency (ms)
PID	0.784	High	< 5
NMPC	0.350	Low	50–100
DDPG	0.505	Medium	10–20
Decision Transformer	1.717	High	30–50

The results must be understood in terms of multi-objective performance: in rigid-body systems, success is synonymous with minimising cross-track error. In articulated systems, hitch stability takes precedence over exact path following. An agent that maintains $|\gamma| < 5^\circ$ throughout a reverse task is more successful than one that achieves lower CTE but triggers jackknife events.

Table 7.2: Experiment matrix: tasks, observations, and success criteria

Experiment	Task	Observation	Success Metric
Path fidelity (forward)	Path following	BEV image + vector state	Mean Abs. CTE < 0.2 m
Dynamic obstacle avoidance	Avoidance	BEV image + vector state	Zero collisions
Lane detection	Perception	Camera image	F1-score of lane mask
Straight-line backing	Reverse	BEV image + vector state	$ \gamma < 5^\circ$ over 20 m
90-degree docking	Reverse	BEV image + vector state	Trailer in bay; no collision
Jackknife recovery	Stability	Vector state	Recover to $ \gamma < 5^\circ$ in 5 m
Sensor degradation	Robustness	Noisy BEV image	CTE increase < 50% vs. clean

Chapter 8

Conclusion

8.1 Summary of Findings

This thesis presented a reinforcement learning framework for articulated autonomous vehicle control, building on the foundational work from ME 780 (DDPG and Decision Transformer for single-vehicle path following) and ECE 682 (LQR control of the vehicle bicycle model with tractor-trailer kinematics). The key contributions demonstrated are:

- a combined dynamic tractor and kinematic trailer model implemented in a custom Gymnasium simulation environment, extending the single-vehicle bicycle model from prior work;
- a vision-based observation space incorporating BEV images rendered from the Pygame simulation, enabling spatial perception of lane boundaries and obstacles;
- a TD3 reinforcement learning agent trained on lane-following and obstacle avoidance tasks with a multiplicative reward function penalising cross-track error, heading error, and hitch articulation angle;
- comparison of TD3 performance against PID and MPC classical control baselines and the prior ME 780 DDPG results.

The integration of visual perception via BEV imagery extends the state-only approach of the ME 780 work. By processing rendered BEV images through a CNN feature extractor, the agent gains spatial context about lane geometry and obstacle positions that vector state observations alone cannot provide.

8.2 Limitations

The study is subject to the following constraints:

- Experiments are conducted entirely within the custom 2D Pygame/Gymnasium simulation environment; performance in higher-fidelity simulators or on physical hardware has not yet been characterised.
- The 2D simulation approximates the vehicle dynamics but does not capture all effects present in a real tractor-trailer system, such as load transfer, tire saturation at high slip angles, or hitch joint compliance.
- The current framework assumes a fixed payload; hitch dynamics change significantly with varying trailer mass.
- Obstacle avoidance is limited to static obstacles; dynamic obstacle handling is not addressed.

8.3 Future Work

Several directions remain open for future investigation:

- **Higher-fidelity simulation:** Validating the trained policies in a physics-based 3D simulator (e.g., Gazebo) with a detailed URDF model of the tractor-trailer system, including realistic tire and hitch dynamics.
- **Hardware deployment:** Transferring trained policies to a physical robotic platform, requiring sim-to-real transfer techniques such as domain randomisation and system identification.
- **Decision Transformer for tractor-trailer:** Applying the DT approach from ME 780 to the articulated vehicle setting, with training data generated from the TD3 expert policy and augmented with high-error recovery trajectories.
- **Multi-trailer systems:** Extending the kinematic model and RL policy to double or B-train configurations.
- **Online adaptation:** Updating the policy in real time to compensate for varying payloads and road conditions.
- **Control Barrier Function overlays:** Implementing CBFs as a hard safety layer that overrides the RL agent when a collision or jackknife is imminent.
- **Dynamic obstacles:** Extending the obstacle avoidance environment to include moving obstacles and integrating predictive collision avoidance.

Appendix A

Simulation Software Architecture

Appendix B

Full Hyperparameter Tables

Table B.1: TD3 Hyperparameters

Hyperparameter	Value
Learning rate	
Batch size	
Replay buffer size	
Policy noise	
Target update rate	
Timesteps	200,000

Appendix C

Additional Plots

Appendix D

Reward Function Derivation

Bibliography